

MML minor #9

Введение в TensorFlow

TensorFlow DL framework

- We will use it in Jupyter Notebook with Python 3 kernel

```
import numpy as np  
import tensorflow as tf
```

- We will overview Python API for TensorFlow 1.2+
- APIs in other languages exist: Java, C++, Go
 - Python API is at present the most complete and the easiest to use
 - https://www.tensorflow.org/api_docs/



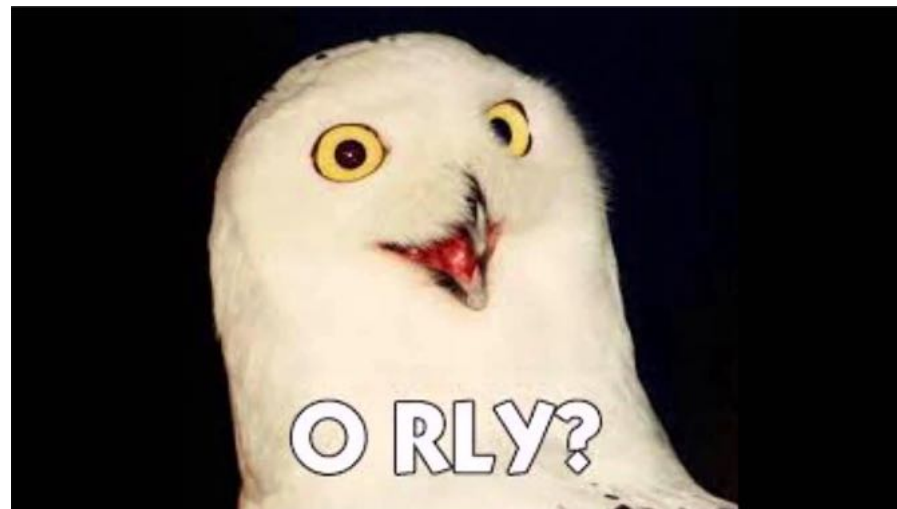
What is TensorFlow?

1. A tool to describe computational graphs
 - The foundation of computation in TensorFlow is the **Graph** object. This holds a network of nodes, each representing one **operation**, connected to each other as inputs and outputs.
2. A runtime for execution of these graphs
 - On CPU, GPU, TPU, ...
 - On one node or in distributed mode



Why this name

- **Input** to any operation will be a collection of **tensors** (multi-dimensional arrays)
- **Output** will be a collection of tensors as well.
- We will have a graph of operations, each of which transforms tensors into another tensors, so it's a kind of a flow of tensors



How the input looks like

- **Placeholder**

- This is placeholder for a tensor, which will be fed during graph execution (e.g. input features)
- `x = tf.placeholder(tf.float32, (None, 10))`

- **Variable**

- This is a tensor with some value that is updated during execution (e.g. weights matrix in MLP)
- `w = tf.get_variable("w", shape=(10, 20), dtype=tf.float32)`
- `w = tf.Variable(tf.random_uniform((10, 20)), name="w")`

- **Constant**

- This is a tensor with constant value, that cannot be changed
- `c = tf.constant(np.ones((4, 4)))`

Operation example

- Matrix product:

```
x = tf.placeholder(tf.float32, (None, 10))
w = tf.Variable(tf.random_uniform((10, 20)), name="w")
z = x @ w
# z = tf.matmul(x, w)
print(z)
```

- Output:

Tensor("matmul:0", shape=(?, 20), dtype=float32)

- We don't do any computations here, we just **define** the graph!

Computational graph

- TensorFlow creates a **default graph** after importing
 - All the operations will go there by default
 - You can get it with `tf.get_default_graph()`, which returns an instance of `tf.Graph`.
- You can **create your own** graph variable and define operations there:

```
g = tf.Graph()  
with g.as_default():  
    pass
```

- You can **clear** the default graph like this:

```
tf.reset_default_graph()
```

Jupyter Notebook cells

- If you run this cell **3 times**:

```
x = tf.placeholder(tf.float32, (None, 10))
```

- This is what you get in your **default graph**:

- Using `tf.get_default_graph().get_operations()`

```
[<tf.Operation 'Placeholder' type=Placeholder>,  
 <tf.Operation 'Placeholder_1' type=Placeholder>,  
 <tf.Operation 'Placeholder_2' type=Placeholder>]
```

- **Your graph is cluttered!**

- Clear your graph with `tf.reset_default_graph()`

Operations and tensors

- Every **node** in our graph is an **operation**:

```
x = tf.placeholder(tf.float32, (None, 10), name="x")
```

- Listing nodes with `tf.get_default_graph().get_operations()`:

```
[<tf.Operation 'x' type=Placeholder>]
```

- How to get **output tensors** of operation:

- `tf.get_default_graph().get_operations()[0].outputs`

- Output: `[<tf.Tensor 'x:0' shape=(?, 10) dtype=float32>]`

Running a graph

- A **tf.Session** object encapsulates the environment in which **tf.Operation** objects are executed, and **tf.Tensor** objects are evaluated.
- Create a session: `s = tf.InteractiveSession()`
- Defining a graph:
`a = tf.constant(5.0)`
`b = tf.constant(6.0)`
`c = a * b`
- Running a graph: `print(c)` *# here just looking at the type*
`print(s.run(c))` *# that's how you run the graph*
- Output:
`Tensor("mul:0", shape=(), dtype=float32)`
`30.0`

Running a graph

- Operations are written in C++ and executed on CPU or GPU.
- `tf.Session` owns necessary resources to execute your graph, such as `tf.Variable`, that occupy RAM.
- It is important to release these resources when they are no longer required with `tf.Session.close()`

Initialization of variables

- A variable has an initial value:
 - Tensor: `tf.Variable(tf.random_uniform((10, 20)), name="w")`
 - Initializer: `tf.get_variable("w", shape=(10, 20), dtype=tf.float32)`
- You need to run some code to **compute that initial value** in graph execution environment
- This is done with a call in your session `s`:

```
s.run(tf.global_variables_initializer())
```
- Without it you will get “Attempting to use uninitialized value” errors

Example

- Definition:

```
tf.reset_default_graph()
a = tf.constant(np.ones((2, 2), dtype=np.float32))
b = tf.Variable(tf.ones((2, 2)))
c = a @ b
```

- Running attempt:

```
s = tf.InteractiveSession()
s.run(c)
```

Output: **“Attempting to use uninitialized value”** error

- Running properly:

```
s.run(tf.global_variables_initializer())
s.run(c)
```

Output: array([[2., 2.], [2., 2.]], dtype=float32)

Feeding placeholder values

- Definition:

```
tf.reset_default_graph()
a = tf.placeholder(np.float32, (2, 2))
b = tf.Variable(tf.ones((2, 2)))
c = a @ b
```

- Running attempt:

```
s = tf.InteractiveSession()
s.run(tf.global_variables_initializer())
s.run(c)
```

Output: **“You must feed a value for placeholder tensor”** error

- Running properly:

```
s.run(tf.global_variables_initializer())
s.run(c, feed_dict={a: np.ones((2, 2))})
```

Output: array([[2., 2.], [2., 2.]], dtype=float32)

Summary

- TensorFlow: defining and running computational graphs
- Nodes of a graph are operations, that convert a collection of tensors into another collection of tensors
- In Python API you **define** the graph, you **don't execute** it along the way
 - In 1.5+ the latter mode is supported: eager execution
- You create a **session** to execute your graph (fast C++ code on CPU or GPU)
- Session **owns** all the resources (tensors eat RAM)

Optimizers in TensorFlow

- Let's define **f** as a square of variable **x**:

```
import numpy as np
```

```
import tensorflow as tf
```

```
tf.reset_default_graph()
```

```
x = tf.get_variable("x", shape=(), dtype=tf.float32)
```

```
f = x ** 2
```

- Let's say we want to *minimize* the value of **f** w.r.t **x**:

```
optimizer = tf.train.GradientDescentOptimizer(0.1)
```

```
step = optimizer.minimize(f, var_list=[x])
```


Trainable variables

- You don't have to specify all the optimized variables:

```
step = optimizer.minimize(f, var_list=[x])  
step = optimizer.minimize(f)
```

- Because all variables are **trainable** by default:

```
x = tf.get_variable("x", shape=(), dtype=tf.float32)  
x = tf.get_variable("x", shape=(), dtype=tf.float32, trainable=True)
```

- You can get all of them:

```
tf.trainable_variables()
```

- Output:

```
[<tf.Variable 'x:0' shape=() dtype=float32_ref>]
```

Making gradient descent steps

- Now we need to create a **session** and **initialize** variables:

```
s = tf.InteractiveSession()  
s.run(tf.global_variables_initializer())
```

- We are ready to make 10 **gradient descent** steps:

```
for i in range(10):  
    _, curr_x, curr_f = s.run([step, x, f])  
    print(curr_x, curr_f)
```

- Output:

```
0.448929 0.314901  
0.359143 0.201537  
...  
0.0753177 0.00886368  
0.0602542 0.00567276
```

← *GD step is already
applied to x*

Logging with tf.Print

- We can **evaluate** tensors and print them like this:

```
for i in range(10):  
    _, curr_x, curr_f = s.run([step, x, f])  
    print(curr_x, curr_f)
```

- Or we can pass our tensor of interest through **tf.Print**:

```
...  
f = x ** 2  
f = tf.Print(f, [x, f], "x, f:")  
...  
for i in range(10):  
    s.run([step, f])
```

```
x, f:[1.5879565][2.521606]  
x, f:[1.2703652][1.6138278]  
...  
x, f:[0.26641488][0.070976891]  
x, f:[0.2131319][0.04542521]
```

*This is Jupyter Server stdout.
Not visible in the Notebook!*

Logging with TensorBoard


- We can add so-called **summaries**:

```
tf.summary.scalar('curr_x', x)
tf.summary.scalar('curr_f', f)
summaries = tf.summary.merge_all()
```

- This is how we log these summaries:

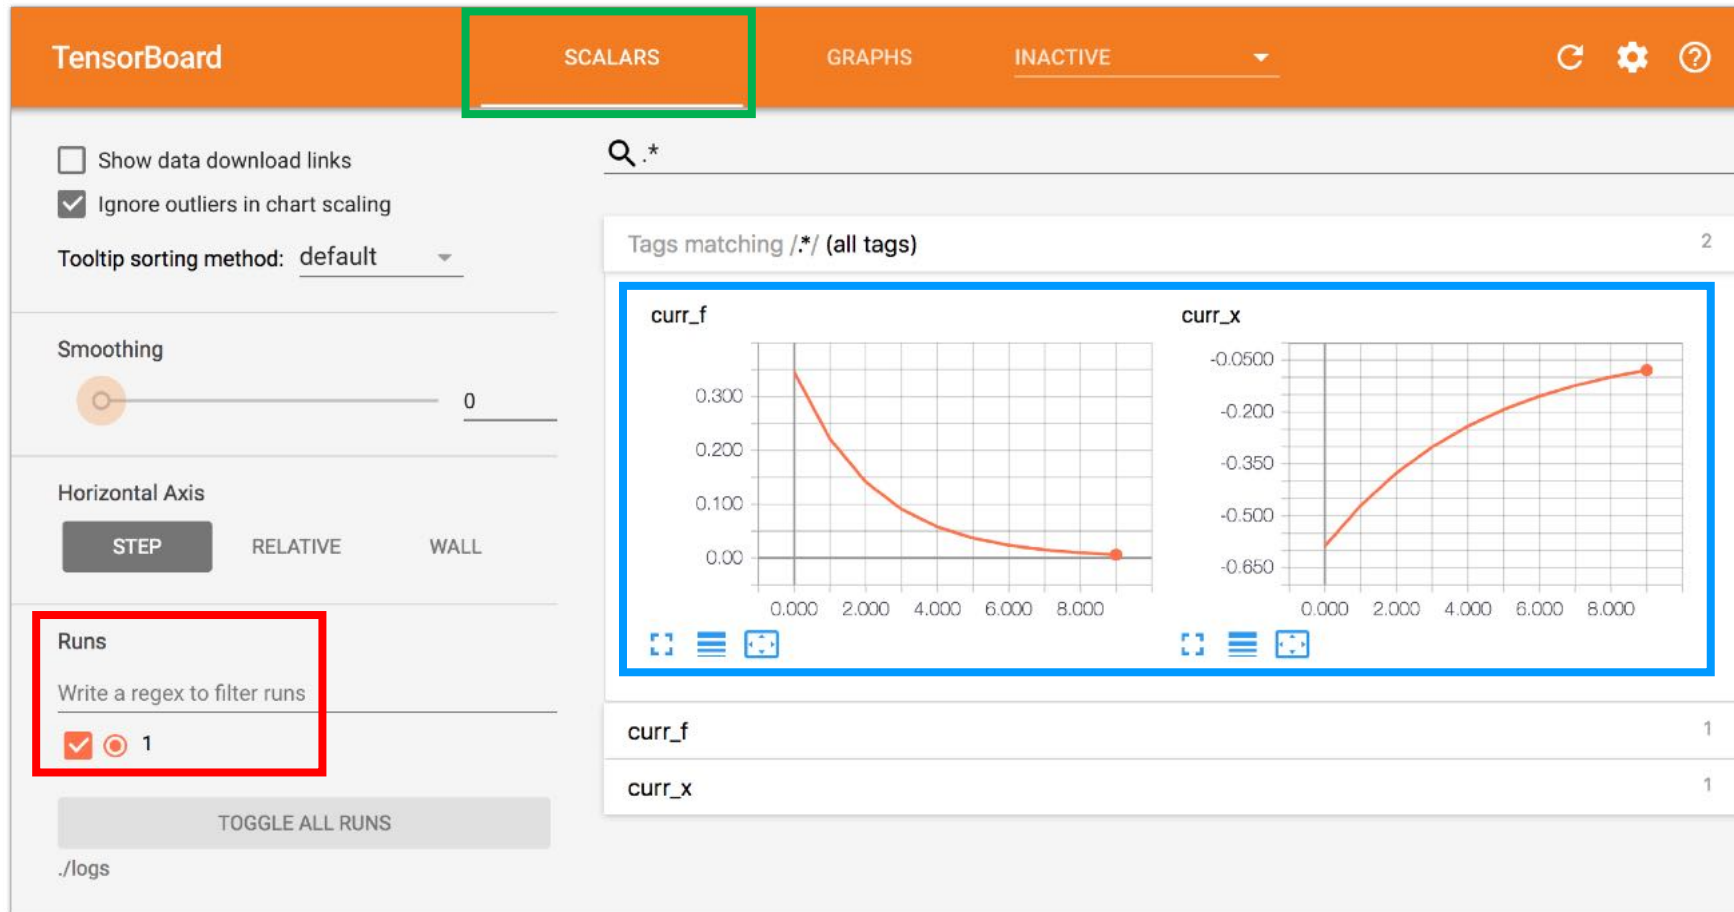
```
s = tf.InteractiveSession()
summary_writer = tf.summary.FileWriter("logs/1", s.graph)
s.run(tf.global_variables_initializer())
for i in range(10):
    _, curr_summaries = s.run([step, summaries])
    summary_writer.add_summary(curr_summaries, i)
    summary_writer.flush()
```

Run number



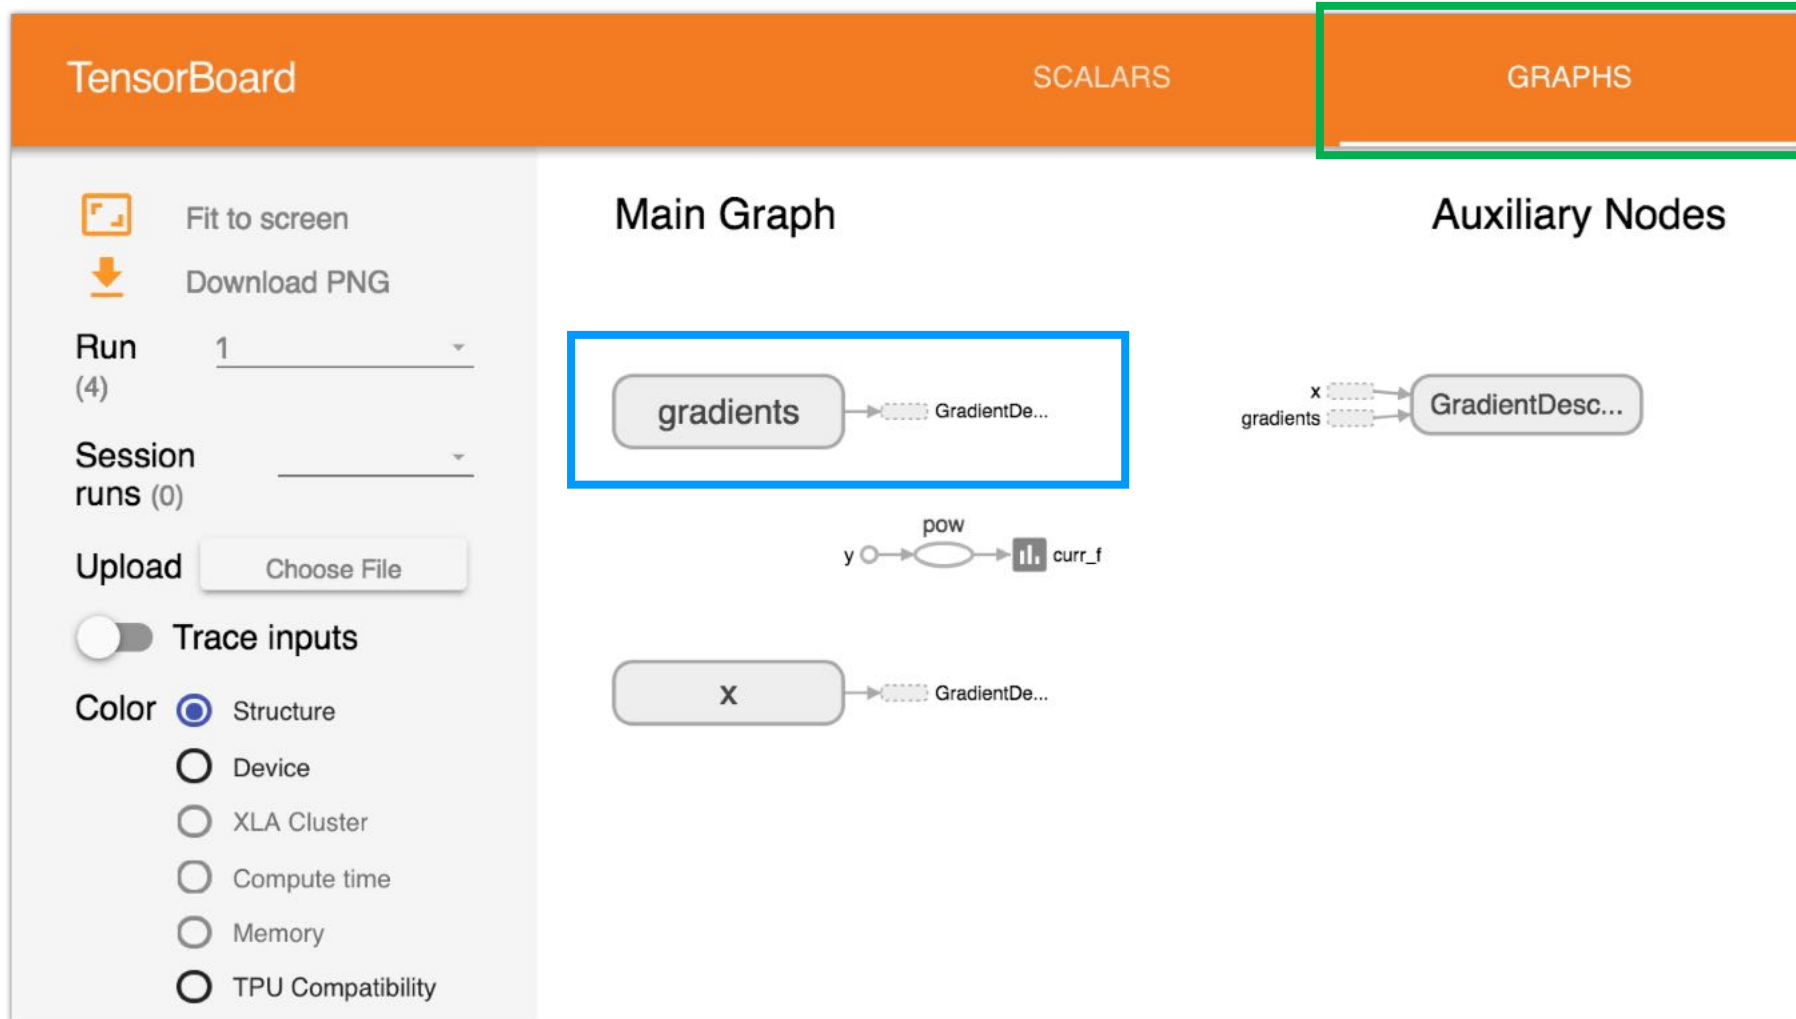
Launching TensorBoard

- Now you can launch TensorBoard via bash:
`tensorboard --logdir=./logs`
- And open **`http://127.0.0.1:6006`** in your browser.



Visualizing graph in TensorBoard

- You can see that **gradients computation** is a part of our graph:



Solving a linear regression

- Let's generate a model dataset:

$N = 1000$

$D = 3$

$x = \text{np.random.random}((N, D))$

$w = \text{np.random.random}((D, 1))$

$y = x @ w + \text{np.random.randn}(N, 1) * 0.20$

Solving a linear regression

- We will need **placeholders** for input data:

```
tf.reset_default_graph()
features = tf.placeholder(tf.float32, shape=(None, D))
target = tf.placeholder(tf.float32, shape=(None, 1))
```

- This is how we make **predictions**:

```
weights = tf.get_variable("w", shape=(D, 1), dtype=tf.float32)
predictions = features @ weights
```

- And define our **loss**:

```
loss = tf.reduce_mean((target - predictions) ** 2)
```

- And **optimizer**:


```
optimizer = tf.train.GradientDescentOptimizer(0.1)
step = optimizer.minimize(loss)
```


Solving a linear regression

- Gradient descent:

```
s = tf.InteractiveSession()
s.run(tf.global_variables_initializer())
for i in range(300):
    _, curr_loss, curr_weights = s.run(
        [step, loss, weights], feed_dict={features: x, target: y})
    if i % 50 == 0:
        print(curr_loss)
```

Filling placeholders



- Ground truth weights:

[0.11649134,0.82753164,0.46924019]

- Found weights:

[0.13715988,0.79555332,0.47024861]

Model checkpoints

- We can save variables' state with **tf.train.Saver**:

```
s = tf.InteractiveSession()
saver = tf.train.Saver(tf.trainable_variables())
s.run(tf.global_variables_initializer())
for i in range(300):
    _, curr_loss, curr_weights = s.run(
        [step, loss, weights], feed_dict={features: x, target: y})
    if i % 50 == 0:
        saver.save(s, "logs/2/model.ckpt", global_step=i)
    print(curr_loss)
```

Model checkpoints

- We can **list** last checkpoints:

```
saver.last_checkpoints  
  
['logs/2/model.ckpt-50', 'logs/2/model.ckpt-100',  
'logs/2/model.ckpt-150', 'logs/2/model.ckpt-200',  
'logs/2/model.ckpt-250']
```

- We can **restore** a previous checkpoint like this:

```
saver.restore(s, "logs/2/model.ckpt-50")
```

- Only variables' values are restored, which means that you need to define a graph in *the same* way **before restoring** a checkpoint.

Summary

- TensorFlow has **built-in optimizers** that do back-propagation automatically.
- TensorBoard provides **tools for visualizing** your training progress.
- TensorFlow allows you to **checkpoint** your graph to restore its state later (*you need to define it in exactly the same way though*)

Ссылки

- <https://www.oreilly.com/learning/hello-tensorflow>
- https://www.tensorflow.org/guide/low_level_intro